

NARRATIVE DETECTION MVP - IMPLEMENTATION ROADMAP

1. Setup & Tech Stack

What to do

Set up the development environment and infrastructure for the entire system.

How

Step 1: Install Core Dependencies

```
# Docker & Docker Compose
curl -fsSL https://get.docker.com | sh

# Ollama (Local LLM)
curl -fsSL https://ollama.com/install.sh | sh
ollama pull llama3.2:3b
```

Step 2: Initialize Project Structure

```
mkdir narrative-detection-mvp && cd narrative-detection-mvp
mkdir -p app/{models,detection,services,api,templates}
mkdir -p data/{raw_json,reports}
mkdir -p tests
```

Step 3: Create Configuration Files

- `docker-compose.yml` - Orchestrate PostgreSQL, Redis, API, Worker
- `requirements.txt` - Python dependencies
- `.env` - Environment variables
- `Dockerfile` - Container image for API and workers

Use (Tech Stack Explained)

Backend Core:

- **Python 3.11+:** Primary language
 - Why: Rich ML/NLP libraries, fast development
 - Fast enough for real-time processing
- **FastAPI:** Web framework for REST API
 - Why: Auto-generated docs, async support, type validation
 - Alternative considered: Flask (simpler but no async)

Data Storage:

- **PostgreSQL 16 + pgvector:** Primary database
 - Why: Stores tweets, users, narratives AND vector embeddings
 - pgvector extension: Enables similarity search without separate vector DB
 - Alternative considered: MongoDB (but we need ACID + vectors)
- **Redis:** Message queue + cache
 - Why: Queues your JSON data between scraper and processor
 - Also caches URL expansions and frequently accessed data
 - Alternative considered: RabbitMQ (more features but heavier)

Processing:

- **asyncio + threading:** Async workers
 - Why: Process multiple tweets concurrently
 - No need for Celery (overkill for MVP)

Frontend:

- **HTMX + Jinja2:** Dynamic UI without JavaScript framework
 - Why: Simpler than React, faster to build
 - Chart.js for visualizations
 - Alternative considered: React (too complex for MVP)

Infrastructure:

- **Docker Compose:** Local orchestration
 - Why: One command starts entire stack
 - Easy team collaboration (same environment)

Task Assignment

- **Person A (Usmaan & Muwafaq):** Docker setup, database schema, environment config
- **Person D (Nadeem):** Verify installations, document setup process

Deliverables

-
- All services running via `docker-compose up`
 - Database accessible at `localhost:5432`
 - Redis accessible at `localhost:6379`
 - API returns health check at `localhost:8000/health`

2. Twitter Data Collector

What to build

Integration layer that consumes your existing 3-layer JSON output and feeds it into the system.

How

Architecture:

Your Scraper → JSON Files → Ingest Script → Redis Streams → Database

Step 1: Redis Streams Setup Create three streams for your three layers:

```
import redis

r = redis.Redis(host='localhost', port=6379)

# Add MICRO layer tweet
r.xadd('tweets:micro', {
    'layer': 'MICRO',
    'data': json.dumps(micro_json)
})

# Add MINUTE layer batch
r.xadd('tweets:minute', {
    'layer': 'MINUTE',
    'data': json.dumps(minute_json)
})

# Add HOURLY layer profile
r.xadd('tweets:hourly', {
    'layer': 'HOURLY',
    'data': json.dumps(hourly_json)
})
```

Step 2: File Watcher (Optional)

```
# watches data/raw_json/ folder
# automatically ingests new JSON files
from watchdog.observers import Observer
from watchdog.events import FileSystemEventHandler

class JSONHandler(FileSystemEventHandler):
    def on_created(self, event):
        if event.src_path.endswith('.json'):
            ingest_json_file(event.src_path)
```

Step 3: Manual Ingest Script

```
python scripts/ingest.py data/raw_json/micro_layer.json
python scripts/ingest.py data/raw_json/minute_layer.json
python scripts/ingest.py data/raw_json/hourly_layer.json
```

Use

Input Processing:

- **JSON parsing:** Python `json` module
 - Validates structure against expected schema
 - Handles malformed data gracefully
- **Redis Streams:** Message queue
 - **Why streams instead of pub/sub:** Persistence + consumer groups
 - Each layer gets its own stream
 - Workers consume from streams without blocking scrapers

Data Flow:

1. Your scraper outputs JSON to `data/raw_json/`
2. Ingest script reads JSON
3. Validates required fields (`tweet_id`, `text`, `timestamp`)
4. Pushes to appropriate Redis stream
5. Returns immediately (non-blocking)

Error Handling:

- Missing fields: Log warning, fill with NULL
- Duplicate `tweet_id`: Skip with log entry
- Malformed JSON: Save to `data/errors/` for manual review

Task Assignment

- **Person A (Usmaan & Muwafaq):** Build ingest script, Redis integration, file watcher

Deliverables

- Script: `python ingest.py <json_file>` works
- Redis streams populated: Check with `redis-cli XLEN tweets:micro`
- Error handling logs issues without crashing

3. Preprocessing & Storage

What to build

Worker process that consumes from Redis, cleans data, and writes to PostgreSQL.

How

Worker Architecture:

```
# worker.py - Runs continuously
import redis
import psycopg2
from preprocessing import clean_tweet

r = redis.Redis()

while True:
    # Read from stream (blocks until data available)
    messages = r.xread({'tweets:micro': '$'}, count=10, block=5000)

    for stream, entries in messages:
        for msg_id, data in entries:
            tweet = json.loads(data['data'])

            # Clean and process
            processed = clean_tweet(tweet)

            # Save to PostgreSQL
            save_to_db(processed)

            # Acknowledge message
            r.xack('tweets:micro', 'worker-group', msg_id)
```

Cleaning Pipeline:

```
def clean_tweet(tweet):
    text = tweet['text_raw']

    # 1. Remove extra whitespace
    text = ' '.join(text.split())

    # 2. Remove/normalize emojis
    text_clean = text.encode('ascii', 'ignore').decode()

    # 3. Extract features
    hashtags = re.findall(r'#\w+', text)
    mentions = re.findall(r '@\w+', text)
    urls = re.findall(r 'https?://[^\\s]+', text)

    # 4. Generate text hash (for duplicate detection)
    normalized = ''.join(c.lower() for c in text if c.isalnum())
    text_hash = hashlib.md5(normalized.encode()).hexdigest()

    return {
        'id': tweet['tweet_id'],
        'text_raw': text,
        'text_clean': text_clean,
        'hashtags': hashtags,
        'mentions': mentions,
        'urls': urls,
        'text_hash': text_hash,
        'created_at': parse_timestamp(tweet['timestamp'])
    }
```

Use

Text Cleaning:

- **Python re module:** Regular expressions for pattern extraction
 - Extract hashtags: #\w+
 - Extract mentions: @\w+
 - Extract URLs: https?:\/\/[^\s]+
- **Unicode handling:** Remove emojis while preserving text
 - Method: `text.encode('ascii', 'ignore').decode()`
 - Alternative: `emoji` library for replacement

Feature Extraction:

- **Hashtags:** Store as PostgreSQL array `TEXT[]`
 - Enables queries: "Find all tweets with #scam"
- **Mentions:** Track interaction network
 - Later used for graph construction
- **URLs:** Critical for coordination detection
 - Many bots share the same malicious link

Duplicate Detection:

- **MD5 text hash:**
 - Normalize: lowercase + remove punctuation
 - Hash the normalized text
 - Same hash = exact duplicate (even if spacing differs)

Database Operations:

- **SQLAlchemy ORM:** Database abstraction
- ```
from sqlalchemy import create_engine
```
- ```
from models import Tweet
```
-
- ```
 tweet = Tweet(
```
- ```
        id=processed['id'],
```
- ```
 text_raw=processed['text_raw'],
```
- ```
        text_clean=processed['text_clean'],
```
- ```
 hashtags=processed['hashtags']
```
- ```
    )
```
- ```
 session.add(tweet)
```
- ```
    session.commit()
```
- **Batch inserts:** Process 100 tweets at once
 - Why: Reduces DB round-trips (10x faster)
 - Use: `session.bulk_insert_mappings()`

Timestamp Handling:

- **dateutil.parser:** Parse various date formats

- from dateutil import parse timestamp = parser.parse("2025-12-03T16:54:22.000Z")
 - Handles ISO8601, RFC3339, Unix timestamps
 - Converts to UTC automatically

Task Assignment

- **Person A (Usmaan & Muwafaq):** Build cleaning pipeline, database models
- **Person B (Omama & Hashir):** Test feature extraction accuracy

Deliverables

- Worker consumes from Redis and writes to PostgreSQL
 - All tweets have `text_clean`, `text_hash`, features extracted
 - Batch processing: 100+ tweets/second throughput
-

4. Narrative Detection (Topics + Spikes)

What to build

Clustering algorithm that groups similar tweets into narratives and detects volume spikes.

How

Two-Step Process:

Step 1: Clustering (Group Similar Tweets)

```
from sentence_transformers import SentenceTransformer
from sklearn.cluster import HDBSCAN
import numpy as np

# Load embedding model
model = SentenceTransformer('all-MiniLM-L6-v2')

# Get recent tweets
tweets = fetch_recent_tweets(hours=6)
texts = [t['text_clean'] for t in tweets]

# Generate embeddings
embeddings = model.encode(texts, show_progress_bar=True)

# Cluster
clusterer = HDBSCAN(
    min_cluster_size=5,          # Need at least 5 tweets
    min_samples=3,               # Density threshold
    metric='euclidean',
    cluster_selection_epsilon=0.5
)
clusters = clusterer.fit_predict(embeddings)
```

```

# Group tweets by cluster
narratives = {}
for idx, cluster_id in enumerate(clusters):
    if cluster_id == -1: # Noise (unclustered)
        continue
    if cluster_id not in narratives:
        narratives[cluster_id] = []
    narratives[cluster_id].append(tweets[idx])

```

Step 2: Spike Detection (Volume Anomalies)

```

def detect_spike(narrative_id):
    # Get tweet timestamps for this narrative
    tweets = get_narrative_tweets(narrative_id)
    timestamps = [t['created_at'] for t in tweets]

    # Calculate baseline (past 24 hours)
    now = datetime.utcnow()
    baseline_start = now - timedelta(hours=24)
    baseline_tweets = [t for t in tweets if t['created_at'] >=
    baseline_start]
    baseline_rate = len(baseline_tweets) / 24 # tweets per hour

    # Current rate (last hour)
    current_start = now - timedelta(hours=1)
    current_tweets = [t for t in tweets if t['created_at'] >=
    current_start]
    current_rate = len(current_tweets)

    # Detect spike (3x threshold)
    velocity = current_rate / max(baseline_rate, 0.1)
    is_spike = velocity >= 3.0

    return {
        'is_spike': is_spike,
        'velocity': velocity,
        'current_rate': current_rate,
        'baseline_rate': baseline_rate
    }

```

Use

Embeddings (Semantic Vectors):

- **sentence-transformers:** Convert text → 384-dimensional vectors
 - Model: all-MiniLM-L6-v2
 - Why this model:
 - Fast (50 tweets/sec on CPU)
 - Good quality (outperforms TF-IDF)
 - Small download (80MB)
 - Alternative: paraphrase-MiniLM-L6-v2 (slower but better quality)
- **Embedding properties:**
 - Similar meaning = similar vectors
 - Captures context (not just keywords)
 - Example: "5G causes cancer" similar to "5G health risks"

Clustering Algorithm:

- **HDBSCAN:** Hierarchical density-based clustering
 - Why not KMeans: We don't know # of narratives in advance
 - Why not DBSCAN: HDBSCAN better handles varying densities
 - Automatically finds number of clusters
 - Labels outliers as noise (cluster_id = -1)
- **Parameters explained:**
 - min_cluster_size=5: Need at least 5 tweets to form narrative
 - min_samples=3: 3 tweets needed in neighborhood for core point
 - cluster_selection_epsilon=0.5: Merge similar clusters

Storage:

- **pgvector:** Store embeddings in PostgreSQL
- ```
Store embedding
tweet.embedding = embeddings[idx].tolist() # Convert numpy to list
session.commit() # Later: Find similar
tweets_similar = session.query(Tweet).order_by(
 Tweet.embedding.cosine_distance(query_embedding)).limit(10).all()
```

## Spike Detection Math:

- **Simple ratio-based:**
  - Velocity = (current hourly rate) / (24hr baseline rate)
  - Spike threshold: 3x (tunable parameter)
  - Example: 30 tweets/hour now vs 10 tweets/hour baseline = 3x = SPIKE

## Temporal Bucketing:

- **pandas groupby:** Group tweets by time windows
- ```
import pandas as pd
df = pd.DataFrame(tweets)
df['time_bucket'] = df['created_at'].dt.floor('10min')
volume = df.groupby('time_bucket').size()
```

Task Assignment

- **Person B (Omama & Hashir):** Implement clustering, test on sample data
- **Person C (Omama & Hashir):** Build spike detection logic

Deliverables

- ✓ Clustering groups similar tweets into narratives
 - ✓ Each narrative has: title, summary, tweet_count, time_range
 - ✓ Spike detector flags 3x+ volume increases
 - ✓ API endpoint: GET /narratives returns list
-

5. Bot Detection (Basic Heuristics)

What to build

Rule-based scoring system that calculates bot probability for each account.

How

Score Calculation:

```
class BotDetector:
    def __init__(self):
        self.weights = {
            'posting_frequency': 0.30,
            'account_age': 0.25,
            'follower_ratio': 0.20,
            'repeat_text': 0.25
        }

    def score_user(self, user):
        score = 0.0

        # Feature 1: Posting frequency
        # Suspicious: >50 posts/day, Bot: >100
        posts_per_day = user['tweet_count'] / max(user['account_age_days'],
1)
        freq_score = min(posts_per_day / 100, 1.0)
        score += freq_score * self.weights['posting_frequency']

        # Feature 2: Account age (newer = suspicious)
        if user['account_age_days'] < 7:
            age_score = 1.0
        elif user['account_age_days'] < 30:
            age_score = 0.7
        elif user['account_age_days'] < 90:
            age_score = 0.3
        else:
            age_score = 0.0
        score += age_score * self.weights['account_age']

        # Feature 3: Follower ratio (anomalies)
        ratio = user['followers'] / max(user['following'], 1)
        if ratio < 0.1 or ratio > 10:
            ratio_score = 0.8
        elif ratio < 0.3 or ratio > 5:
            ratio_score = 0.5
        else:
            ratio_score = 0.0
        score += ratio_score * self.weights['follower_ratio']

        # Feature 4: Repeat text (get from DB)
        repeat_ratio = calculate_repeat_ratio(user['user_id'])
        repeat_score = min(repeat_ratio / 0.5, 1.0)
        score += repeat_score * self.weights['repeat_text']

        # Classify
        if score >= 0.7:
            label = 'BOT'
```

```

        elif score >= 0.4:
            label = 'SUSPICIOUS'
        else:
            label = 'ORGANIC'

    return score, label

def calculate_repeat_ratio(user_id):
    # Get all tweets from user
    tweets = get_user_tweets(user_id)

    # Count unique text hashes
    unique_hashes = set(t['text_hash'] for t in tweets)

    # Ratio of repeated content
    return 1 - (len(unique_hashes) / len(tweets))

```

Use

Feature Engineering:

- **Posts per day:** tweet_count / account_age_days
 - Calculate from HOURLY profile data
 - Bots often have 100+ posts/day
 - Stored in users.posts_per_day column
- **Account age:** (today - account_created).days
 - Parse join date from profile: "Joined August 2011"
 - Use dateutil.parser for flexible parsing
 - New accounts (<30 days) often used for campaigns
- **Follower ratio:** followers / following
 - Normal users: 0.5 - 2.0
 - Bots: Often <0.1 (follow many, few followers) or >10 (bought followers)
 - Handle division by zero: max(following, 1)
- **Repeat text ratio:** SQL query
 - WITH user_tweets AS (
 - SELECT text_hash FROM tweets WHERE user_id = ?
 -)
 - SELECT
 - COUNT(*) as total_tweets,
 - COUNT(DISTINCT text_hash) as unique_tweets,
 - 1 - (COUNT(DISTINCT text_hash)::float / COUNT(*)) as repeat_ratio
 - FROM user_tweets

Scoring Formula:

- **Weighted sum:** Each feature contributes to final score
- **Weights (sum = 1.0):**
 - Posting frequency: 30% (most reliable indicator)
 - Account age: 25% (new accounts risky)
 - Follower ratio: 20% (good signal but can be gamed)
 - Repeat text: 25% (strong bot indicator)

Thresholds:

- Score 0.0-0.4: ORGANIC (normal user)
- Score 0.4-0.7: SUSPICIOUS (manual review)
- Score 0.7-1.0: BOT (high confidence)

Why This Approach:

- **Transparent:** Show users exactly why flagged
- **Tunable:** Adjust weights without retraining
- **Fast:** Pure math (1000 profiles/second)
- **No training data:** Works immediately

Update Schedule:

- Recalculate scores when:
 - New profile data arrives (HOURLY layer)
 - User posts 10+ new tweets
 - Manual request via API

Task Assignment

- **Person C (Omama & Hashir):** Implement bot scoring, test on sample profiles

Deliverables

- Bot scores calculated for all users
 - Scores stored in `users.bot_score` column
 - Labels (ORGANIC/SUSPICIOUS/BOT) assigned
 - API endpoint: `GET /users/{handle}` shows bot score
-

6. Coordinated Attack + Cross-Account Similarity

What to build

Detection system that finds groups of accounts posting identical/similar content in tight time windows.

How

Sliding Window Algorithm:

```
class CoordinationDetector:  
    def __init__(self, time_window_minutes=10, similarity_threshold=0.85):  
        self.time_window = time_window_minutes
```

```

        self.similarity_threshold = similarity_threshold

def detect_coordination(self, tweets):
    clusters = []

    # Group by text hash (exact duplicates)
    hash_groups = {}
    for tweet in tweets:
        h = tweet['text_hash']
        if h not in hash_groups:
            hash_groups[h] = []
        hash_groups[h].append(tweet)

    # Find coordinated groups
    for text_hash, group in hash_groups.items():
        if len(group) < 3: # Need at least 3 accounts
            continue

        # Check time window
        timestamps = sorted([t['created_at'] for t in group])
        time_span = (timestamps[-1] - timestamps[0]).total_seconds() /
60

        if time_span <= self.time_window:
            # Get unique users
            users = list(set(t['user_id'] for t in group))

            if len(users) >= 3: # At least 3 different accounts
                clusters.append({
                    'text_hash': text_hash,
                    'users': users,
                    'tweet_ids': [t['id'] for t in group],
                    'time_span_minutes': time_span,
                    'sample_text': group[0]['text_clean']
                })

    # Also check embedding similarity (catches paraphrased content)
    semantic_clusters = self.find_semantic_similarity(tweets)

    return clusters + semantic_clusters

def find_semantic_similarity(self, tweets):
    """Find similar (not identical) content using embeddings"""
    from sklearn.metrics.pairwise import cosine_similarity

    # Get embeddings
    embeddings = np.array([t['embedding'] for t in tweets])

    # Calculate similarity matrix
    sim_matrix = cosine_similarity(embeddings)

    # Find highly similar groups
    clusters = []
    processed = set()

    for i in range(len(tweets)):
        if i in processed:
            continue

        # Find tweets similar to tweet i

```

```

        similar_idx = np.where(sim_matrix[i] >
self.similarity_threshold)[0]

        if len(similar_idx) >= 3: # At least 3 similar tweets
            similar_tweets = [tweets[j] for j in similar_idx]

            # Check time window
            times = [t['created_at'] for t in similar_tweets]
            time_span = (max(times) - min(times)).total_seconds() / 60

            if time_span <= self.time_window:
                users = list(set(t['user_id'] for t in similar_tweets))

                if len(users) >= 3:
                    clusters.append({
                        'type': 'SEMANTIC',
                        'users': users,
                        'tweet_ids': [t['id'] for t in similar_tweets],
                        'avg_similarity':
float(np.mean(sim_matrix[i][similar_idx])),
                        'time_span_minutes': time_span
                    })
                    processed.update(similar_idx)

    return clusters

```

Use

Exact Duplicate Detection:

- **Text hashing:** MD5 of normalized text
 - Already calculated in preprocessing stage
 - Stored in `tweets.text_hash` column
 - **SQL query:** `SELECT * FROM tweets WHERE text_hash = ? GROUP BY user_id HAVING COUNT(*) >= 3`

Sliding Window:

- **Time-based grouping:** Group tweets within N minutes
 - # PostgreSQL query with time window
 - `SELECT text_hash, array_agg(user_id), array_agg(id)`
 - `FROM tweets`
 - `WHERE created_at > NOW() - INTERVAL '24 hours'`
 - `GROUP BY domain`
 - `HAVING COUNT(DISTINCT user_id) >= 5`
 - `ORDER BY account_count DESC`
 - **Coordination signal:** 5+ accounts linking to same obscure domain = suspicious

Integration with Tweets:

- **Database storage:** Update `tweets.expanded_urls` column
- # After expansion `tweet.expanded_urls = [result['domain'] for result in expanded_results]` `session.commit()`

Rate Limiting:

- Respect domain rate limits
- Use connection pooling
- Batch requests (don't expand URLs one-by-one)

Task Assignment

- **Person A (Usmaan & Muwafaq):** Build URL expansion service
- **Person B (Omama & Hashir):** Maintain suspicious domain list

Deliverables

- URL expander service with Redis caching
 - All shortened URLs resolved to final destinations
 - Suspicious domains flagged
 - Shared domain analysis in coordination detection
-

8. Community Detection (Graph)

What to build

Social network graph that identifies clusters of accounts based on interactions and content similarity.

How

Graph Construction:

```
import networkx as nx

def build_interaction_graph(tweets, users):
    G = nx.Graph()

    # Add nodes (accounts)
    for user in users:
        G.add_node(
            user['user_id'],
            handle=user['handle'],
            bot_score=user['bot_score']
        )

    # Add edges (interactions)
    for tweet in tweets:
        # Retweet edges
        if tweet.get('retweeted_from'):
            G.add_edge(
                tweet['user_id'],
                tweet['retweeted_from'],
                edge_type='RETWEET',
                weight=1.0
```

```

        )

    # Reply edges
    if tweet.get('reply_to_user'):
        G.add_edge(
            tweet['user_id'],
            tweet['reply_to_user'],
            edge_type='REPLY',
            weight=0.8
        )

    # Mention edges
    for mention in tweet.get('mentions', []):
        if mention in users_dict:
            G.add_edge(
                tweet['user_id'],
                users_dict[mention],
                edge_type='MENTION',
                weight=0.5
            )

    # Add similarity edges (content-based)
    similarity_edges = find_similar_content_pairs(tweets)
    for user_a, user_b, similarity in similarity_edges:
        if G.has_edge(user_a, user_b):
            # Strengthen existing edge
            G[user_a][user_b]['weight'] += similarity
        else:
            G.add_edge(user_a, user_b, edge_type='SIMILAR',
                       weight=similarity)

    return G

def find_similar_content_pairs(tweets):
    """Find pairs of users posting similar content"""
    from sklearn.metrics.pairwise import cosine_similarity

    # Group by user
    user_embeddings = {}
    for tweet in tweets:
        uid = tweet['user_id']
        if uid not in user_embeddings:
            user_embeddings[uid] = []
        user_embeddings[uid].append(tweet['embedding'])

    # Average embeddings per user
    user_avg_embeddings = {
        uid: np.mean(embs, axis=0)
        for uid, embs in user_embeddings.items()
    }

    # Find similar pairs
    users = list(user_avg_embeddings.keys())
    embeddings = np.array([user_avg_embeddings[u] for u in users])
    sim_matrix = cosine_similarity(embeddings)

    pairs = []
    for i in range(len(users)):
        for j in range(i+1, len(users)):
            if sim_matrix[i][j] > 0.7: # High similarity
                pairs.append((users[i], users[j], sim_matrix[i][j]))

```

```
    return pairs
```

Community Detection:

```
from networkx.algorithms import community

def detect_communities(G):
    # Louvain algorithm (fast, good quality)
    communities = community.louvain_communities(G, weight='weight')

    # Analyze each community
    results = []
    for idx, comm in enumerate(communities):
        members = list(comm)

        # Calculate community stats
        bot_scores = [G.nodes[n]['bot_score'] for n in members]
        avg_bot_score = np.mean(bot_scores)

        # Internal vs external edges
        internal_edges = G.subgraph(members).number_of_edges()
        external_edges = sum(1 for u in members for v in G.neighbors(u) if
v not in comm)

        # Classify community
        if avg_bot_score > 0.6:
            community_type = 'BOT_CLUSTER'
        elif len(members) > 20 and internal_edges / len(members) > 3:
            community_type = 'COORDINATED_GROUP'
        else:
            community_type = 'ORGANIC'

        results.append({
            'community_id': idx,
            'size': len(members),
            'members': members,
            'avg_bot_score': avg_bot_score,
            'internal_edges': internal_edges,
            'external_edges': external_edges,
            'type': community_type
        })

    return results
```

Use

Graph Library:

- **NetworkX:** Python graph library
 - Why: Easy to use, good for <10K nodes (sufficient for MVP)
 - Alternative: Neo4j (for production scale >100K nodes)
 - In-memory processing (fast for analysis)

Node Types:

- **User nodes:** Each account is a node

- Attributes: handle, bot_score, account_age
- Stored in graph memory, not separate DB

Edge Types (weighted):

- **RETWEET**: weight=1.0 (strong signal of agreement/amplification)
- **REPLY**: weight=0.8 (engagement)
- **MENTION**: weight=0.5 (weaker signal)
- **SIMILAR**: weight=0.0-1.0 (based on content similarity)

Why weighted edges:

- Community detection considers edge weights
- Stronger connections = more likely same community
- Captures both explicit (retweet) and implicit (similar content) relationships

Community Detection Algorithm:

- **Louvain method**: Fast modularity optimization
- ```
from networkx.algorithms import community
```
- ```
# Detect communities
```
- ```
communities = community.louvain_communities(
```
- ```
    G,
```
- ```
 weight='weight', # Use edge weights
```
- ```
    resolution=1.0 # Default (higher = more communities)
```
- ```
)
```
- **How it works:**
  1. Start with each node in its own community
  2. Iteratively merge communities that increase modularity
  3. Modularity: (internal edges - expected random) / total edges
  4. Stops when no more improvement
- **Why Louvain:**
  - Fast: O(n log n) time
  - Good quality: Near-optimal modularity
  - Hierarchical: Can zoom into sub-communities

### Alternative Algorithms:

- **Greedy modularity**: Simpler, almost as good
- **Label propagation**: Faster but lower quality
- **Girvan-Newman**: Slow but interpretable (removes edges)

### Community Classification:

- **BOT\_CLUSTER**: avg\_bot\_score > 0.6
  - High concentration of bot accounts
  - Likely coordinated bot network
- **COORDINATED\_GROUP**: Many internal edges

- High interconnectivity
- Members frequently interact
- **ORGANIC:** Normal user groups
  - Natural clustering around topics

## Graph Metrics:

- **Modularity:** How well-separated are communities
  - Value: -0.5 to 1.0 (higher = better separation)
  - Good: >0.3
- **Density:** edges / possible\_edges
  - Bot networks: Often high density (everyone connects)
  - Organic: Lower density (selective connections)

## Visualization (optional for MVP):

```
import matplotlib.pyplot as plt

pos = nx.spring_layout(G) # Position nodes
colors = [G.nodes[n]['bot_score'] for n in G.nodes()]

nx.draw(G, pos,
 node_color=colors,
 cmap=plt.cm.RdYlGn_r,
 with_labels=False,
 node_size=50)
plt.savefig('graph.png')
```

## Task Assignment

- **Person C (Omama & Hashir):** Build graph construction, community detection
- **Person B (Omama & Hashir):** Test on sample networks, visualize

## Deliverables

- Interaction graph built from tweets
  - Communities detected and classified
  - Bot clusters identified
  - Graph stored in edges table
  - API endpoint: GET /communities lists clusters
- 

## 9. Narrative Origin Identification

### What to build

Timeline analysis that traces each narrative back to its earliest tweets (origin seeds).

## How

### Origin Detection:

```
def find_narrative_origin(narrative_id):
 # Get all tweets in narrative
 tweets = session.query(Tweet).filter(
 Tweet.narrative_id == narrative_id
).order_by(Tweet.created_at).all()

 if not tweets:
 return None

 # First N tweets are origin candidates
 origin_window_minutes = 30
 first_tweet_time = tweets[0].created_at
 cutoff_time = first_tweet_time +
timedelta(minutes=origin_window_minutes)

 origin_seeds = [
 t for t in tweets
 if t.created_at <= cutoff_time
]

 # Analyze origin characteristics
 origin_users = list(set(t.user_id for t in origin_seeds))
 origin_bot_scores = [
 get_bot_score(uid) for uid in origin_users
]

 # Build spread timeline
 timeline = build_spread_timeline(tweets)

 return {
 'first_tweet_id': tweets[0].id,
 'first_tweet_time': tweets[0].created_at,
 'origin_seeds': [
 {
 'tweet_id': t.id,
 'user_handle': t.handle,
 'bot_score': get_bot_score(t.user_id),
 'text': t.text_clean,
 'timestamp': t.created_at
 }
 for t in origin_seeds
],
 'origin_user_count': len(origin_users),
 'origin_bot_ratio': np.mean(origin_bot_scores),
 'spread_timeline': timeline
 }

def build_spread_timeline(tweets):
 """Group tweets into time buckets"""
 import pandas as pd

 df = pd.DataFrame([
 {'timestamp': t.created_at, 'id': t.id}
 for t in tweets
])
```

```

Group by 5-minute buckets
df['time_bucket'] = df['timestamp'].dt.floor('5min')
timeline = df.groupby('time_bucket').size().to_dict()

Convert to list format
return [
 {
 'time': bucket.isoformat(),
 'tweet_count': count
 }
 for bucket, count in sorted(timeline.items())
]

```

## Spread Velocity Analysis:

```

def calculate_spread_metrics(timeline):
 """Analyze how fast narrative spreads"""
 if len(timeline) < 2:
 return {}

 # Time to reach milestones
 cumulative = 0
 milestones = {}

 for bucket in timeline:
 cumulative += bucket['tweet_count']

 if 'first_10' not in milestones and cumulative >= 10:
 milestones['first_10'] = bucket['time']
 if 'first_50' not in milestones and cumulative >= 50:
 milestones['first_50'] = bucket['time']
 if 'first_100' not in milestones and cumulative >= 100:
 milestones['first_100'] = bucket['time']

 # Calculate velocity (tweets per hour)
 first_time = datetime.fromisoformat(timeline[0]['time'])
 last_time = datetime.fromisoformat(timeline[-1]['time'])
 duration_hours = (last_time - first_time).total_seconds() / 3600
 total_tweets = sum(b['tweet_count'] for b in timeline)
 velocity = total_tweets / max(duration_hours, 0.1)

 # Find peak
 peak_bucket = max(timeline, key=lambda b: b['tweet_count'])

 return {
 'total_tweets': total_tweets,
 'duration_hours': round(duration_hours, 2),
 'velocity': round(velocity, 2), # tweets/hour
 'peak_time': peak_bucket['time'],
 'peak_volume': peak_bucket['tweet_count'],
 'milestones': milestones
 }

```

## Use

### Timeline Construction:

- **pandas groupby**: Efficient time bucketing
- import pandas as pd

- df['time\_bucket'] = pd.to\_datetime(df['timestamp']).dt.floor('5min')
- counts = df.groupby('time\_bucket').size()
- **Bucket sizes:**
  - 5 minutes: Detailed view (use for fast-spreading narratives)
  - 10 minutes: Standard view (most narratives)
  - 1 hour: Long-term trends

### Origin Seed Selection:

- **Time window approach:** First 30 minutes
  - Why 30 min: Captures initial propagation
  - Too short (5 min): Might miss early amplifiers
  - Too long (2 hours): Includes secondary spread
- **Alternative: First N tweets:** Take first 5-10 tweets
  - Simpler but less robust
  - Might miss simultaneous origins

### Spread Patterns:

- **Viral (organic):**
  - Slow start → exponential growth → plateau
  - Origin from high-follower accounts
  - Low bot ratio in origin
- **Coordinated (attack):**
  - Sudden burst → sustained volume → drop
  - Multiple simultaneous origins
  - High bot ratio in origin (>0.6)

### Temporal Analysis:

- **Velocity:** tweets per hour
  - Organic viral: 50-200 tweets/hour at peak
  - Coordinated: 500+ tweets/hour burst
- **Milestones:** Time to reach 10, 50, 100 tweets
  - Fast (organic viral): 10 tweets in <15 min
  - Fast (coordinated): 10 tweets in <5 min
  - Slow (organic trend): 10 tweets in hours

### Database Queries:

```
-- Get origin tweets
SELECT * FROM tweets
WHERE narrative_id = ?
ORDER BY created_at ASC
LIMIT 10;

-- Timeline aggregation
SELECT
 date_trunc('minute', created_at, 5) as time_bucket,
```

```

 COUNT(*) as tweet_count
FROM tweets
WHERE narrative_id = ?
GROUP BY time_bucket
ORDER BY time_bucket;

```

### Storage:

- **narratives table:** Store origin metadata
- UPDATE narrativesSET origin\_tweet\_ids = ARRAY[123, 456, 789],  
origin\_handles = ARRAY['@user1', '@user2'], first\_seen = '2025-12-03  
16:54:22' WHERE narrative\_id = ?

### Task Assignment

- **Person B (Omama & Hashir):** Build origin detection, timeline analysis

### Deliverables

- Origin tweets identified for each narrative
  - Spread timeline with 5-minute buckets
  - Velocity metrics calculated
  - API endpoint: GET /narratives/{id}/origin returns origin analysis
- 

## 10. Advisory Countermeasures

### a) Narrative Risk Score

#### What to build

Weighted risk scoring formula that combines multiple signals into single score.

#### How

##### Risk Calculation:

```

class RiskScorer:
 def __init__(self):
 # Configurable weights (sum = 1.0)
 self.weights = {
 'bot_ratio': 0.30,
 'spike_velocity': 0.25,
 'coordination': 0.25,
 'suspicious_urls': 0.20
 }

 def calculate_risk(self, narrative):
 score = 0.0
 details = {}

```

```

1. Bot ratio (0-1, higher = more bots)
bot_ratio = narrative['bot_ratio']
bot_contribution = bot_ratio * self.weights['bot_ratio']
score += bot_contribution
details['bot_ratio'] = {
 'value': bot_ratio,
 'contribution': bot_contribution,
 'interpretation': self.interpret_bot_ratio(bot_ratio)
}

2. Spike velocity (normalize to 0-1)
velocity = narrative['spike_velocity']
5x spike = max risk
velocity_norm = min((velocity - 1) / 4, 1.0)
velocity_contribution = velocity_norm *
self.weights['spike_velocity']
score += velocity_contribution
details['spike_velocity'] = {
 'value': velocity,
 'normalized': velocity_norm,
 'contribution': velocity_contribution
}

3. Coordination score (0-1)
coord = narrative['coordination_score']
coord_contribution = coord * self.weights['coordination']
score += coord_contribution
details['coordination'] = {
 'value': coord,
 'contribution': coord_contribution
}

4. Suspicious URLs (normalize)
url_count = narrative['suspicious_url_count']
5+ URLs = max risk
url_norm = min(url_count / 5, 1.0)
url_contribution = url_norm * self.weights['suspicious_urls']
score += url_contribution
details['suspicious_urls'] = {
 'count': url_count,
 'normalized': url_norm,
 'contribution': url_contribution
}

Classify level
if score >= 0.7:
 level = 'HIGH'
 urgency = 'CRITICAL'
elif score >= 0.4:
 level = 'MEDIUM'
 urgency = 'MODERATE'
else:
 level = 'LOW'
 urgency = 'ROUTINE'

return {
 'risk_score': round(score, 3),
 'risk_level': level,
 'urgency': urgency,
 'breakdown': details
}

```

```

def interpret_bot_ratio(self, ratio):
 if ratio >= 0.7:
 return "SEVERE: Highly automated campaign"
 elif ratio >= 0.4:
 return "MODERATE: Significant bot activity"
 else:
 return "LOW: Mostly organic accounts"

```

## Use

### Scoring Components:

#### 1. Bot Ratio (30% weight)

- **Calculation:** bot\_accounts / total\_accounts
- **Source:** From bot detection (Section 5)
- **Query:**
- ```
SELECT COUNT(*) FILTER (WHERE u.bot_score >= 0.7) / COUNT(*)::float as bot_ratio
FROM tweets t
JOIN users u ON t.user_id = u.user_id
WHERE t.narrative_id = ?
```

2. Spike Velocity (25% weight)

- **Calculation:** current_rate / baseline_rate
- **Source:** From spike detection (Section 4)
- **Normalization:** (velocity - 1) / 4
 - 1x = 0.0 (no spike)
 - 5x = 1.0 (max risk)
 - 10x = still 1.0 (capped)

3. Coordination Score (25% weight)

- **Calculation:** Percentage of tweets in coordinated clusters
- **Source:** From coordination detection (Section 6)
- **Query:**
- ```
SELECT COUNT(*) FILTER (WHERE t.id = ANY (SELECT unnest(tweet_ids) FROM coordination_clusters)) / COUNT(*)::float as coord_score
FROM tweets t
WHERE t.narrative_id = ?
```

#### 4. Suspicious URLs (20% weight)

- **Calculation:** Count of flagged domains
- **Source:** From URL expansion (Section 7)
- **Normalization:** min(count / 5, 1.0)
  - 0 URLs = 0.0
  - 5+ URLs = 1.0

### Threshold Calibration:

- **LOW (0.0-0.4):** Monitor, no immediate action

- **MEDIUM (0.4-0.7):** Investigate, prepare response
- **HIGH (0.7-1.0):** Immediate action required

### Why These Weights:

- Bot ratio (30%): Strong indicator of artificial campaigns
- Spike velocity (25%): Rapid spread often indicates coordination
- Coordination (25%): Direct evidence of organized activity
- URLs (20%): Lower weight (sometimes organic sharing)

### Tunability:

- Weights stored in config file
- Easy to adjust without code changes
- Can A/B test different weight combinations

## b) Response Timing Recommendation

### What to build

Rule-based system that recommends when to respond.

### How

#### Timing Logic:

```
class ResponseAdvisor:
 def recommend_timing(self, narrative):
 risk_level = narrative['risk_level']
 velocity = narrative['spike_velocity']
 volume = narrative['tweet_count']

 # Decision tree
 if risk_level == 'HIGH':
 if velocity >= 3.0:
 return {
 'timing': 'IMMEDIATE',
 'timeframe': '< 30 minutes',
 'rationale': 'High-risk fast-spreading narrative requires immediate containment',
 'priority': 'P0'
 }
 else:
 return {
 'timing': 'URGENT',
 'timeframe': '< 2 hours',
 'rationale': 'High-risk but slower spread allows brief preparation',
 'priority': 'P1'
 }
 elif risk_level == 'MEDIUM':
 if velocity >= 2.0:
 return {
```

```

 'timing': 'DELAY',
 'timeframe': '2-4 hours',
 'rationale': 'Moderate risk, gather more data before
responding',
 'priority': 'P2'
 }
else:
 return {
 'timing': 'MONITOR',
 'timeframe': '6-12 hours',
 'rationale': 'Watch for escalation before committing
response',
 'priority': 'P3'
 }

else: # LOW risk
 return {
 'timing': 'MONITOR',
 'timeframe': '24 hours',
 'rationale': 'Low risk, continue monitoring without
response',
 'priority': 'P4'
 }

```

## Use

### Decision Matrix:

| Risk Level | Velocity           | Timing    | Timeframe  |
|------------|--------------------|-----------|------------|
| HIGH       | Fast ( $\geq 3x$ ) | IMMEDIATE | <30 min    |
| HIGH       | Normal ( $<3x$ )   | URGENT    | <2 hours   |
| MEDIUM     | Fast ( $\geq 2x$ ) | DELAY     | 2-4 hours  |
| MEDIUM     | Normal             | MONITOR   | 6-12 hours |
| LOW        | Any                | MONITOR   | 24+ hours  |

### Rationale:

- **IMMEDIATE:** Viral misinformation spreading rapidly
  - Example: "Bank XYZ collapsing" false rumor
  - Every minute counts to stop panic
- **URGENT:** Serious but not viral yet
  - Example: Coordinated bot attack starting
  - Time to prepare quality response
- **DELAY:** Worth responding but not urgent
  - Example: Moderate criticism gaining traction
  - Better to respond thoughtfully
- **MONITOR:** Watch but don't engage yet

- Example: Low-volume complaint
- Responding might amplify ("Streisand effect")

### **Additional Factors:**

- Time of day: Immediate response harder at 3am
- Day of week: Weekends vs weekdays
- Media involvement: Journalists asking = more urgent

## c) Context-Specific Reply Strategy

### What to build

Classification system that determines appropriate response tone and approach.

### How

#### **Strategy Selection:**

```
class ReplyStrategySelector:
 def select_strategy(self, narrative):
 # Classify narrative type
 narrative_type = self.classify_narrative(narrative)
 bot_ratio = narrative['bot_ratio']
 volume = narrative['tweet_count']

 strategies = {
 'SCAM': {
 'tone': 'FIRM',
 'approach': 'Call out fraud directly with evidence',
 'template': 'WARNING: This is a known scam. [Evidence link]',
 'include_legal': True
 },
 'PANIC': {
 'tone': 'CALM',
 'approach': 'Reassure with facts and authoritative sources',
 'template': 'We understand concerns. Here are the facts: [...]',
 'include_expert': True
 },
 'DEFAMATION': {
 'tone': 'EVIDENCE_BASED',
 'approach': 'Correct with verifiable data and sources',
 'template': 'The facts: [Data]. Sources: [Links]',
 'include_sources': True
 },
 'CRITICISM': {
 'tone': 'POLITE',
 'approach': 'Acknowledge concern and provide context',
 'template': 'We hear you. Here\'s what we\'re doing: [...]',
 'include_action': True
 },
 }
```

```

 'COORDINATED_BOT': {
 'tone': 'TRANSPARENT',
 'approach': 'Expose coordination evidence',
 'template': 'We\'ve detected coordinated inauthentic
activity. [Details]',
 'include_evidence': True
 }
 }

 strategy = strategies.get(narrative_type, strategies['CRITICISM'])

 # Adjust for bot ratio
 if bot_ratio > 0.6:
 strategy['tone'] = 'TRANSPARENT'
 strategy['call_out_bots'] = True

 return strategy

def classify_narrative(self, narrative):
 """Classify narrative based on content and patterns"""
 keywords = narrative.get('top_keywords', [])
 urls = narrative.get('top_domains', [])
 bot_ratio = narrative['bot_ratio']

 # Rule-based classification
 scam_indicators = ['invest', 'guarantee', 'double your', 'limited
time']
 panic_indicators = ['collapse', 'crisis', 'shutdown', 'failing']
 defamation_indicators = ['fraud', 'illegal', 'scandal', 'corrupt']

 text_sample = ' '.join(keywords).lower()

 if any(ind in text_sample for ind in scam_indicators):
 return 'SCAM'
 elif any(ind in text_sample for ind in panic_indicators):
 return 'PANIC'
 elif any(ind in text_sample for ind in defamation_indicators):
 return 'DEFAMATION'
 elif bot_ratio > 0.6:
 return 'COORDINATED_BOT'
 else:
 return 'CRITICISM'

```

## LLM Integration (Ollama):

```

import ollama

def generate_reply(narrative, strategy):
 prompt = f"""You are a social media response specialist. Generate a
brief, professional reply to address this narrative.

Narrative Summary: {narrative['summary']}
Tone: {strategy['tone']}
Approach: {strategy['approach']}

Key facts to include:
- Origin: {narrative['first_seen']}
- Volume: {narrative['tweet_count']} tweets
- Bot involvement: {narrative['bot_ratio']*100:.1f}%

```

```
Generate a reply that:
1. Matches the {strategy['tone']} tone
2. Is under 280 characters (Twitter limit)
3. {strategy['approach']}
```

Reply:"""

```
response = ollama.generate(
 model='llama3.2:3b',
 prompt=prompt,
 options={
 'temperature': 0.7,
 'max_tokens': 150
 }
)

suggested_reply = response['response'].strip()

return {
 'suggested_reply': suggested_reply,
 'tone': strategy['tone'],
 'approach': strategy['approach'],
 'requires_approval': True, # Always require human review
 'template_used': strategy.get('template')
}
```

## Use

### Narrative Classification:

- **Keyword matching:** Simple but effective
  - Scam: "invest", "guarantee", "limited"
  - Panic: "collapse", "crisis", "failing"
  - Use any(keyword in text for keyword in indicators)

### Tone Selection:

- **FIRM:** For scams, fraud
  - Direct language, no hedging
  - Include legal warnings
- **CALM:** For panic/fear narratives
  - Reassuring language
  - Cite authoritative sources
- **EVIDENCE\_BASED:** For false claims
  - Lead with data
  - Link to verifiable sources
- **POLITE:** For legitimate criticism
  - Acknowledge validity
  - Explain actions taken
- **TRANSPARENT:** For bot attacks
  - Show detection evidence
  - Call out inauthenticity

## LLM (Ollama) Integration:

- **Model:** llama3.2:3b
  - Why: Fast (< 3 seconds), runs locally
  - Alternative: llama3:8b (better quality, slower)
- **Prompt engineering:**
  - Structured prompt with clear instructions
  - Include narrative context and tone guidance
  - Character limit constraint (Twitter)
- **Temperature:** 0.7
  - Not too creative (0.9+) = might hallucinate
  - Not too boring (0.1) = repetitive

## Human-in-Loop:

- **Always require approval:** LLMs can make mistakes
- Show suggested reply in dashboard
- Allow editing before posting
- Log all approved/rejected suggestions

## d) Evidence Summary Packet

### What to build

HTML/Markdown report generator that bundles all evidence.

### How

#### Report Generation:

```
from jinja2 import Template
import matplotlib.pyplot as plt

def generate_evidence_report(narrative_id):
 # Gather all data
 narrative = get_narrative(narrative_id)
 origin = get_origin_analysis(narrative_id)
 coordination = get_coordination_clusters(narrative_id)
 risk = calculate_risk(narrative)
 strategy = select_strategy(narrative)

 # Generate timeline chart
 timeline_chart = create_timeline_chart(narrative)

 # Build report data
 report_data = {
 'narrative': narrative,
 'origin': origin,
 'coordination': coordination,
 'risk': risk,
 'strategy': strategy,
 'timeline_chart': timeline_chart,
```

```

 'generated_at': datetime.utcnow().isoformat(),
 'report_id': f"RPT-{narrative_id}-{int(time.time())}"
 }

 # Render HTML
 template = Template(REPORT_TEMPLATE)
 html = template.render(**report_data)

 # Save
 report_path = f"data/reports/{report_data['report_id']}.html"
 with open(report_path, 'w') as f:
 f.write(html)

 # Also '1 hour'
 GROUP BY text_hash
 HAVING COUNT(DISTINCT user_id) >= 3
 AND MAX(created_at) - MIN(created_at) < INTERVAL '10 minutes'

```

## Semantic Similarity:

- **Cosine similarity:** Measure angle between embedding vectors
  - Value 0-1 (1 = identical, 0 = opposite)
  - Threshold 0.85 catches paraphrasing
  - Example: "5G causes cancer" vs "5G health dangers" = 0.89 similarity
- **pgvector query:**
- # Find tweets similar to target
- similar = session.query(Tweet).filter(
 • Tweet.embedding.cosine\_distance(target\_embedding) < 0.15 # 1-
 0.85
 • ).all()

## Coordination Signals:

1. **Same text + same time** = Strong signal (likely coordinated)
2. **Similar text + same time** = Moderate signal (might be organic trend)
3. **Same text + different times** = Weak signal (could be copy/pasta)

## Storage:

- **coordination\_clusters table:** Store detected groups
- INSERT INTO coordination\_clusters ( user\_ids, tweet\_ids, shared\_text\_hash, similarity\_score, time\_window\_minutes, cluster\_size) VALUES ( ARRAY[123, 456, 789], ARRAY[111, 222, 333], 'abc123def', 0.92, 8, 3)

## Real-time vs Batch:

- **Real-time:** Check each incoming tweet against last 1 hour
- **Batch:** Hourly job scans last 24 hours for patterns
- Why both: Real-time for alerts, batch for deep analysis

## Task Assignment

- **Person C (Omama & Hashir):** Implement coordination detection
- **Person B (Omama & Hashir):** Test semantic similarity accuracy

## Deliverables

- Detects groups of 3+ accounts posting same content within 10 minutes
  - Catches paraphrased content via embedding similarity
  - Stores coordination clusters in database
  - API endpoint: GET /coordination lists clusters
- 

## 7. Expanded URLs

### What to build

URL unshortening service that follows redirects and identifies shared malicious domains.

### How

#### URL Expansion Logic:

```
import httpx
from urllib.parse import urlparse
import time

class URLExpander:
 def __init__(self):
 self.cache = {} # Cache expanded URLs
 self.suspicious_domains = set() # Known bad domains

 async def expand_url(self, short_url, timeout=5):
 # Check cache first
 if short_url in self.cache:
 return self.cache[short_url]

 try:
 async with httpx.AsyncClient(follow_redirects=True,
 timeout=timeout) as client:
 response = await client.head(short_url)
 final_url = str(response.url)

 # Extract domain
 domain = urlparse(final_url).netloc

 # Store in cache
 self.cache[short_url] = {
 'final_url': final_url,
 'domain': domain,
 'status_code': response.status_code,
 'expanded_at': time.time()
 }

 except httpx.ConnectError:
 self.suspicious_domains.add(short_url)
```

```

 }

 return self.cache[short_url]

 except Exception as e:
 # URL unreachable or timeout
 return {
 'final_url': short_url,
 'domain': urlparse(short_url).netloc,
 'error': str(e)
 }

def is_suspicious_domain(self, domain):
 """Check against known malicious domains"""
 # Check local blacklist
 if domain in self.suspicious_domains:
 return True

 # Check heuristics
 suspicious_patterns = [
 'bit.ly', # URL shorteners (not inherently bad, but used in
campaigns)
 'tinyurl.com',
 '.tk', # Free TLDs often used for scams
 '.ml',
 '.ga'
]

 return any(pattern in domain for pattern in suspicious_patterns)

```

## Batch Processing:

```

async def process_urls_batch(tweets):
 expander = URLExpander()

 # Extract all URLs
 all_urls = []
 for tweet in tweets:
 all_urls.extend(tweet.get('urls', []))

 # Expand concurrently
 tasks = [expander.expand_url(url) for url in all_urls]
 expanded = await asyncio.gather(*tasks)

 # Find shared domains
 domain_counts = {}
 for result in expanded:
 domain = result['domain']
 domain_counts[domain] = domain_counts.get(domain, 0) + 1

 # Flag domains shared by many accounts
 suspicious = {
 domain: count
 for domain, count in domain_counts.items()
 if count >= 5 # 5+ accounts sharing same domain
 }

 return suspicious

```

## Use

### HTTP Client:

- **httpx**: Async HTTP library
  - Why not `requests`: `httpx` supports async (much faster for batch)
  - `follow_redirects=True`: Automatically follows 301/302
  - `timeout=5`: Don't wait forever for dead links

### Caching Strategy:

- **Redis cache**: Store expanded URLs
- `import redis`
- `r = redis.Redis()`
- 
- `# Cache for 7 days`
- `r.setex(f'url:{short_url}', 604800, json.dumps(expanded_url))`
- 
- `# Check cache`
- `cached = r.get(f'url:{short_url}')`
- `if cached:`
- `return json.loads(cached)`
- **Why cache**: Same URLs appear in many tweets
  - Reduces HTTP requests by 80%+
  - Faster response times

### Domain Extraction:

- **urllib.parse**: Parse URL components
- `from urllib.parse import urlparse`  
`url = "https://example.com/path?param=value"`  
`parsed = urlparse(url)`  
`domain = parsed.netloc # "example.com"`

### Suspicious Domain Detection:

- **Blacklist approach**: Maintain list of known bad domains
  - Update periodically from threat intelligence feeds
  - Store in `suspicious_domains` set ( $O(1)$  lookup)
- **Heuristic approach**: Pattern matching
  - Free TLDs: .tk, .ml, .ga, .cf (often used for scams)
  - URL shorteners: bit.ly, tinyurl (not bad, but used in campaigns)
  - Newly registered domains: Check domain age via WHOIS

### Shared Domain Analysis:

- **Aggregation**: Count how many accounts share each domain
- `SELECT unnest(expanded_urls) as domain, COUNT(DISTINCT user_id) as account_count`  
FROM tweets  
WHERE created\_at > NOW() - INTERVAL